

# Databases

BCS1510

Dr. Katharina Schneider &

Dr. Tony Garnock-Jones\*



Week 4 - Lecture 1




EPD150 MSM Conference Hall

Adapted with gratitude from the original lectures by Dr. Ashish Sai

\*CC both of us when emailing please!



Maastricht University



# What have we looked at so far?

- Database and SQL Fundamentals
  - Overview of **databases**, **DBMS evolution**, **data models**, and SQL essentials.
  - Introduction to **relational algebra** (see also lab 1) and **SQL operations** (table creation + basic data manipulation)
- Advanced SQL Techniques and Operations
  - **Advanced query techniques** (handling **NULL values**, using **complex conditions**, and combining data across multiple tables through different **JOIN operations**).
- Subqueries and Operators
  - Introduction to the role of **subqueries** in SQL and their usage in SELECT, FROM and WHERE clauses
  - Detailed discussion on **operators** such as IN, EXISTS, ANY and ALL
- Aggregation, Grouping and Subquery Applications
  - Use of **SQL aggregate functions** (SUM, AVG, COUNT, MIN, MAX)
  - Understanding grouping data with the **GROUP BY** clause and applying the **HAVING** clause
  - **Practical applications** of subqueries in complex SQL series

# Learning Objectives

- **Understand Set Operations in SQL**
  - Learn to use set operations like UNION, INTERSECT and MINUS to combine, intersect, or subtract results from different queries
  - Understand the impact of these operations on the result sets
- **Exploring Built-In Functions**
  - Gain familiarity with SQL's built-in functions (including their application in data manipulation and query refinement).
  - Emphasize the usage of standard functions like SUM, AVG, MIN, MAX, and COUNT
- **Mastering Data Modification Statements:**
  - Develop skills to effectively use INSERT, DELETE and UPDATE statements in SQL to modify data within databases
  - Learn how to handle default values and use subqueries within these statements to manipulate large datasets
- **Applying Advanced SQL Techniques**
  - Apply advanced SQL techniques in practical scenarios, such as using subqueries in data modification, handling complex conditions in deletion, and managing database updates to enforce business rules or constraints
- **Using SQL in Host Language Environment**
  - Understand the three tier architecture
  - Learn how to process SQL statement with JDBC

# Set Operators

UNION, UNION ALL, INTERSECT and  
MINUS



# Set Operators

- Set operators and joins both combine data from multiple tables
- Remember JOINS:
  - Combine columns from multiple tables
  - Can join tables with different columns
  - Aim: bring related data together
- So what do Set Operators do?

# Set Operators

- Combine results of entire queries (rows)
- Unite, intersect, or subtract only the results of subqueries with the same schema, i.e., same number of attributes, in the same order and with the same data types
- Set Operators in SQL remove duplicated rows in the results by default (Use UNION ALL if you want to suppress the deletion of duplicates)

# Set Operator: UNION

- <subquery> UNION <subquery>
- Q: Using tables PCs, Laptops and Printers, find the list of prices of all products. List the model number and price for each product

```
SELECT model, price FROM PCs
UNION
SELECT model, price FROM Laptops
UNION
SELECT model, price FROM Printers;
```

- Note that the three subqueries above have the same schema

# Set Operator: MINUS

- <subquery> MINUS <subquery>
- Q: Using Products(maker, model, type), find the makers who make at least one laptop model but no PC models

```
SELECT maker FROM Products WHERE type = 'laptop'  
MINUS  
SELECT maker FROM Products WHERE type = 'pc';
```

- Note that MINUS is not implemented in MySQL
- How can we do it instead?



# MINUS - Alternative

- Q: Using Products(maker, model, type), find the makers who make at least one laptop model but no PC models

```
SELECT DISTINCT maker
FROM Products
WHERE type = 'laptop' AND maker NOT IN
  (SELECT maker FROM Products WHERE type = 'pc');
```

# Set Operator: INTERSECT

- <subquery> INTERSECT <subquery>
- Q: Using Products, PCs and Laptops, find the makers who make at least one PC model with price above 500 and at least one laptop model also with price above 500.

```
SELECT maker FROM Products NATURAL JOIN PCs WHERE price > 500  
INTERSECT  
SELECT maker FROM Products NATURAL JOIN Laptops WHERE price > 500;
```

- Note that INTERSECT is not implemented in MySQL either
- How can we do it instead?

# INTERSECT - Alternative

- Q: Using Products, PCs and Laptops, find the makers who make at least one PC model with price above 500 and at least one laptop model also with price above 500.

```
SELECT DISTINCT maker FROM Products NATURAL JOIN PCs
WHERE price > 500 AND maker IN
  (SELECT maker FROM Products NATURAL JOIN Laptops
   WHERE price > 500);
```

# Built-In Functions



# Built-In Functions

- SQL standard defines a set of standard SQL built-in functions
- Only some aggregate functions (SUM, AVG, MIN, MAX, COUNT) are covered in this course
- but you are highly encouraged to use a wider range of built-in functions in your project

# Usage of Built-in Functions

- See aggregate functions in last lecture

# Built-in Functions: Some references

- Standard built-in functions in SQL92
  - <http://db.apache.org/derby/docs/10.4/ref/rrefsqj29026.html>
- Built-in functions in MySQL:
  - <http://dev.mysql.com/doc/refman/5.7/en/functions.html>
- Built-in functions in Oracle:
  - [https://docs.oracle.com/cd/B28359\\_01/server.111/b28286/functions001.htm#SQLRF51174](https://docs.oracle.com/cd/B28359_01/server.111/b28286/functions001.htm#SQLRF51174)
- Built-in functions in the Microsoft SQL Server:
  - <http://db.apache.org/derby/docs/10.4/ref/rrefsqj29026.html>

# Data Modification Statements

INSERT, DELETE and UPDATE





# Data Modification Statements

- A modification command does not return a result but changes the database in some way
- We already saw the basic data modification statements in lecture 2
- Remember:

# INSERT

# Inserting a Tuple

- To insert a single new tuple:

```
INSERT INTO table_name(column1, column2, column3, ...)  
VALUES (value_column1, value_column2, value_column3);
```

- Example:

```
INSERT INTO customers(customer_id, firstname,  
                        lastname, city)  
VALUES ('0000000000', 'John', 'Doe', 'Maastricht');
```

# Default Values

- In a CREATE TABLE statement, we can follow an attribute by DEFAULT and a value

```
CREATE TABLE students (  
    id INT primary key,  
    name VARCHAR(100),  
    grade INT DEFAULT 10);
```

- When an inserted tuple has no value for that attribute, the default will be used

# Inserting a Tuple

- Do we really need to specify the attributes in addition to the table?
- No, not necessarily
- But it is helpful if
  - We forgot the standard order of attributes for the relation
  - We don't have values for all attributes and we want the system to fill in missing components with NULL or a default value

# Inserting a Tuple Without Specifying Attributes

- To insert a single new tuple:

```
INSERT INTO table_name  
VALUES (<list of values>);
```

- Example

```
INSERT INTO customers  
VALUES ('00000000001', 'John', 'Doe', 'Maastricht',  
      'Paul-Henri-Spaaklaan 1', NULL);
```

# Inserting Many Tuples from a Subquery

- We may insert the entire result of a query into a relation using:

```
INSERT INTO table_name  
(<subquery>) ;
```

# Inserting Many Tuples from a Subquery

- Create a copy of Products(maker, model, type) called Products\_copy

```
CREATE TABLE Products_copy(  
    maker CHAR(1),  
    model CHAR(4) Primary key,  
    type VARCHAR(8) DEFAULT 'pc');
```

```
INSERT INTO Products_copy  
(SELECT * FROM Products);
```



# DELETE



# Deleting Rows

- To delete a row/multiple rows:

```
DELETE FROM table_name  
WHERE condition;
```

- Example:

```
DELETE FROM Products_copy  
WHERE type = 'printer' AND maker = 'A';
```

- Without condition (WHERE), all tuples in the relation are deleted

# Deleting Tuples: Example

- Delete from Products\_copy(maker, model, type) all models for which there is another mode by the same maker
- Suggestions?

# Deleting Tuples: Example

- Delete from Products\_copy(maker, model, type) all models for which there is another mode by the same maker

```
DELETE FROM Products_copy p
WHERE EXISTS ( SELECT model FROM products_copy WHERE
               maker = p.maker AND model <> p.model);
```

# Deleting some tuples – Semantics

- Delete from Products(maker, model, type) all models for which there is another model by the same maker:
- Suppose maker H makes only products 3006 and 3007
- We go through the tuples one by one and check whether there is another model for the current maker
- Suppose we come to the tuple of product 3006 first
- So there is another model by the same maker and we delete the tuple of product 3006
- What happens to the 3007 tuple?

# Deleting some tuples – Semantics

- Answer: we delete the 3007 tuple as well
- The reason is that the deletion proceeds in two stages:
  - Mark all tuples for which the WHERE condition is satisfied
  - Delete the marked tuples

# UPDATE

# UPDATE

- To change certain attributes in certain tuples of a relation

```
UPDATE <relation>  
SET <list of attribute assignments>  
WHERE <condition on tuples>
```

- Example:

```
UPDATE PCs  
SET price = 1999.0  
WHERE model='1001';
```



# UPDATE Several Tuples

- Make 1999.0 the maximum price for PCs

```
UPDATE PCs  
SET price = 1999.0  
WHERE price > 1999.0;
```


# Short Break

- **10 minutes**



# Tic Tac Toe Game and Database Systems

- Imagine two players play a tic tac toe game together, one using a computer, one using a mobile phone




# Tic Tac Toe Game and Database Systems

- What may be stored in the database?
  - Move records
  - Players
  - Game states
  - ...
- Example:

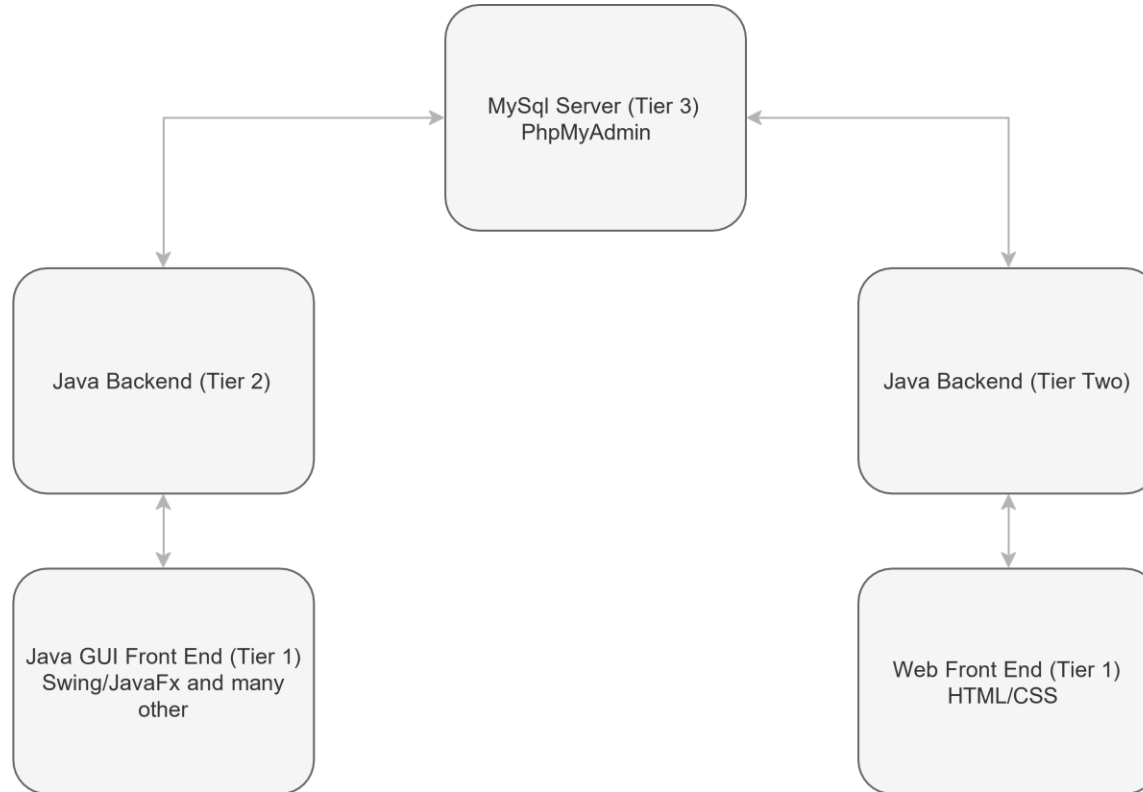
| Move_id | Game_id | Player_id | Position | Symbol | Move_time           |
|---------|---------|-----------|----------|--------|---------------------|
| 1       | 1       | 23        | 0        | X      | 2025-04-15 16:00:01 |
| 2       | 1       | 45        | 1        | O      | 2025-04-15 16:00:10 |

# Architecture Realization

- A common environment for using a database has three tiers (three-tier client-server architecture)
- First Tier – Client (front-end)
  - User Interface
- Second Tier – Application Server (back-end)
  - Business logic
  - Data processing logic
- Third Tier – Database server
  - Data validation
  - Database access



# Three-Tier Architecture Tic Tac Toe




# Environments, Connections, Queries

- The database is, in many DB-access languages an **environment**
- Database servers maintain some number of **connections**, so app servers can ask queries or perform modifications
- The app server issues **statements**: queries and modifications, usually



# Environments, Connections, Queries



# JDBC Java Database Connectivity

<https://docs.oracle.com/javase/tutorial/jdbc/basics/index.html>



# Processing SQL Statements with JDBC

- In general, to process any SQL statement with JDBC, you follow these steps
  - Establishing a connection
  - Create a statement
  - Execute the query
  - Process the ResultSet object
  - Close the connection

# Establish a connection

# Establish a Connection

```
import java.sql.*;

public class main {
    private static String dbUrl = "jdbc:mysql://localhost:3306/pcshop";
    private static String dbUsername = "root";
    private static String dbPassword = "...";

    public static void main(String[] args) {
        try {
            Connection myCon = DriverManager.getConnection(dbUrl,
                dbUsername, dbPassword) ;
        }
        ...
    }
}
```

# Create a Statement

# Create a statement

- JDBC provides classes for three different kinds of statements
  - **Statement**: Used to implement simple SQL statements with no parameters
  - **PreparedStatement**: (extends Statement) Used for precompiling SQL statements that might contain input parameters
  - **CallableStatement**: (extends PreparedStatement) Used to execute stored procedures that may contain both input and output parameters

# Create a Statement

```
Connection conn = DriverManager.getConnection(...); // Establish conn  
  
Statement stmt = conn.createStatement(); // Create the statement  
String sql = "SELECT * FROM Products"; // Create query string
```

- <https://docs.oracle.com/javase/tutorial/jdbc/basics/processingstatements.html>



# Create a Statement

```
Connection conn = DriverManager.getConnection(...); // Establish conn

PreparedStatement pstmt = conn.prepareStatement("SELECT * FROM customers
                                                WHERE customer_id = ?");
pstmt.setString(1, "0000000001");
```

- Placeholder for parameters: ?
- <https://docs.oracle.com/javase/tutorial/jdbc/basics/prepared.html>

# When to use Statement when PreparedStatement?

- Static query with no user input: Statement is okay
- Query with user input of variables: Always use PreparedStatement (SQL injection)
- Repeated execution with different values: PreparedStatement

# SQL injection (Real-World Attack Scenario)

- Scenario:
  - Imagine you have a login form and someone enters
  - Username: admin' --
  - Password: anything

# SQL injection - Statement

```
Connection conn = DriverManager.getConnection(...);
```

```
Statement stmt = conn.createStatement();
```

```
String sql = "SELECT * FROM Users  
WHERE username = 'admin' --' AND password = 'anything'";
```

- -- is a comment in SQL, so the right part gets ignored
- Result: login as an admin without knowing the password

# SQL injection avoided - PreparedStatement

```
Connection conn = DriverManager.getConnection(...);
```

```
String sql = "Select * FROM Users WHERE username = ? AND password = ?";
```

```
PreparedStatement pstmt = conn.createStatement(sql);
```

```
pstmt.setString(1,admin);
```

```
pstmt.setString(2,password);
```

# Execute a Query

# Executing Statements - Updates vs Queries

- JDBC distinguishes queries from modifications, which it calls “updates”
- Statement and PreparedStatement each have methods **executeQuery** and **executeUpdate**
  - Statements: `executeQuery(sql)` – with argument (the query or modification to be executed)
  - PreparedStatements: `executeQuery()` – without argument

# Example: Query + Statement

```
Connection conn = DriverManager.getConnection(...); // Establish conn

Statement stmt = conn.createStatement(); // Create the statement
String sql = "SELECT * FROM Products"; // Create query string
ResultSet rs = stmt.executeQuery(sql); // Execute the query here
```



# Example: Query + PreparedStatement

```
Connection conn = DriverManager.getConnection(...);

String sql = "Select * FROM Users WHERE username = ? AND password = ?";
PreparedStatement pstmt = conn.createStatement(sql);
pstmt.setString(1,admin);
pstmt.setString(2,password);
ResultSet rs = pstmt.executeQuery(); // Execute the query here
```

# Process the ResultSet Object

# Accessing the ResultSet

- An object of type ResultSet is something like a cursor
- Method next() advances the “cursor” to the next tuple
- The first time next() is applied, it gets the first tuple
- If there are no more tuples, next() returns the values false

# Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet
- Method `getX(i)` (where X is some type, and i is the attribute/column number) returns the value of that attribute
- The value must have type X

# Example: Accessing Components

- Q: Give me the total sales for each model.
- SQL query:

```
SELECT model, SUM(quantity)
FROM pcshop.Sales
GROUP BY model;
```

- Java-Code:

```
String sql = "Select model, SUM(quantity) FROM pcshop.sales GROUP BY model";
PreparedStatement pstmt = conn.createStatement(sql);
ResultSet totalSales = pstmt.executeQuery();
```

```
while(totalsales.next()){
String m = totalsales.getString(1); //gets the value in column 1
int tq = totalsales.getFloat(2); //gets the value in column 2
}
```

# Contact details

- If you have any issues throughout the course, contact us via email:
- Tony: [tony.garnock-jones@maastrichtuniversity.nl](mailto:tony.garnock-jones@maastrichtuniversity.nl)
- Katharina: [k.schneider@maastrichtuniversity.nl](mailto:k.schneider@maastrichtuniversity.nl)
- Or send a message on Discord
- (<https://discord.gg/Be2KSF8QG6>)



## Exercises from last lecture

- Q4: Find the customer who bought the most products
- Q5: Find customers who have made purchases on more than one day

# Questions?





# This was my last lecture in that course

See you hopefully around

